

Writing TCP-based Servers



Chris Brown

In This Module ...


Client-side and
server-side
operations

Key data structures
Key sockets system calls

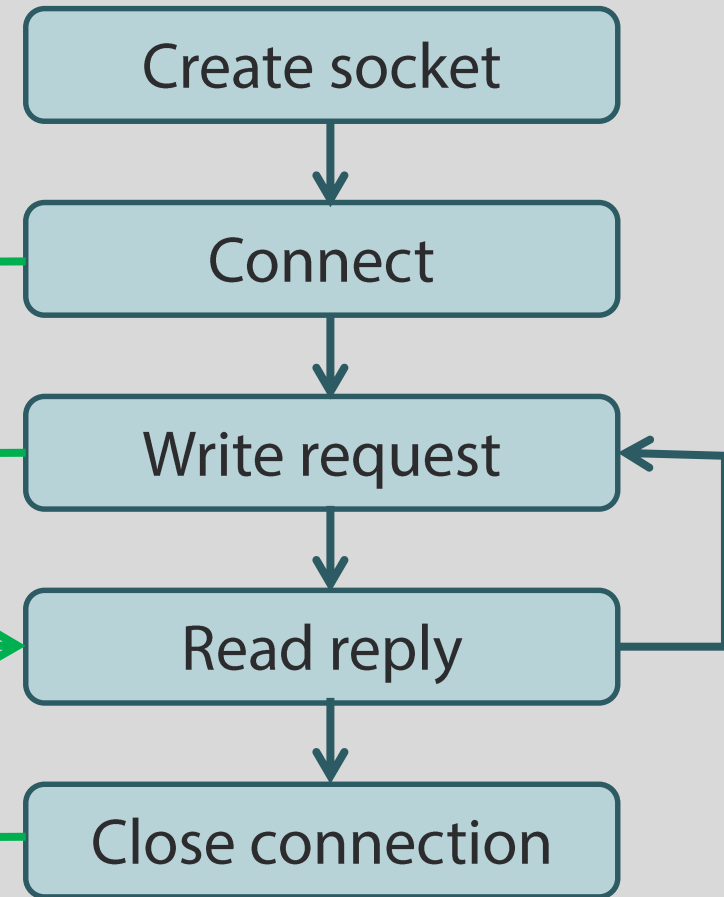
Python as an
alternative language

Demonstration:
A 'rot13' server

Server



Client



The sockaddr_in Structure

```
struct in_addr {  
    in_addr_t  s_addr;  
}
```

← 32-bit IP Address


```
struct sockaddr_in {  
    sa_family_t    sin_family;  
    in_port_t      sin_port;  
    struct in_addr sin_addr;  
    unsigned char  __pad[..];  
}
```

← Address family (AF_INET)

← IPV4 address

← Pad to size of generic sockaddr struct

Little-endian vs. Big-endian



← Little-Endian

← Big-Endian



Network Byte Order

Socket addresses must be in network byte order (big-endian)





Most significant byte stored at lowest memory address, sent across network first

Macros convert to/from the machine's internal byte order.

Host byte order

Network byte order



Creating a Socket

Transport type:
SOCK_STREAM
SOCK_DGRAM



```
sock = socket(domain, type, protocol)
```



Returns an integer
socket descriptor
or -1 on error



The address family. One of:
AF_UNIX
AF_INET
AF_INET6



Usually 0, system
selects protocol
based on domain
and type

Setting the Local Address

```
#define SERVER_PORT      1067
struct sockaddr_in      server;

server.sin_family        = AF_INET;
server.sin_addr.saddr    = htonl(INADDR_ANY)
server.sin_port          = htons(SERVER_PORT);

bind(sock, (struct sockaddr *)&server, sizeof server);
```


Waiting for Business

```
struct sockaddr_in client;  
int fd, client_len;  
  
listen(sock, 5); ← Set up a queue for pending connections  
  
client_len = sizeof(client);  
fd = accept(sock, (struct sockaddr *)&client, &client_len);
```

↑
Connection descriptor

↑
Rendezvous descriptor

↑
Client's endpoint address
returned here.

Talking to the Client

- The connection descriptor returned from `accept()` looks like an open file
 - Supports `read()` and `write()` system calls

```
unsigned char buf[1024];  
int count;  
while ((count = read(fd, buf, 1024)) > 0)  
{  
    rot13(buf, count);  
    write(fd, buf, count);  
}
```

← Pre-allocated buffer

↑ Connection descriptor

Demonstration: The rot13 Server



Doing It in Python



- Python's `sockets` module exposes the traditional sockets API
- Python language features hide some of the messier stuff
 - Automatic buffer allocation on receiving
 - Implicit buffer length management on sending
 - Optional arguments to methods
 - Passing and returning tuples
 - Exception handling replaces error checking
- A socket is a *class*
 - Methods `bind()`, `listen()`, `accept()` ... expose the API

Doing It in Python



```
s = socket(AF_INET, SOCK_STREAM)
s.bind(('', 1068))
s.listen(5)
```

Empty address string
means INADDR_ANY

```
while True:
    client, addr = s.accept()
    rot13_service(client)
    client.close()
```

Tuple assignment:
(connection descriptor,
client endpoint address)

Demonstration: Python Server



Moving Forward ...



In this module:

Sequence of operations

Key data structures

Key system calls

The rot13 server

Doing it in Python

Coming up in the next module:

The client side of TCP/IP