

Writing TCP-based Clients



Chris Brown

In This Module ...


A TCP client for our
rot13 server

Finding the service
(the old-fashioned way)

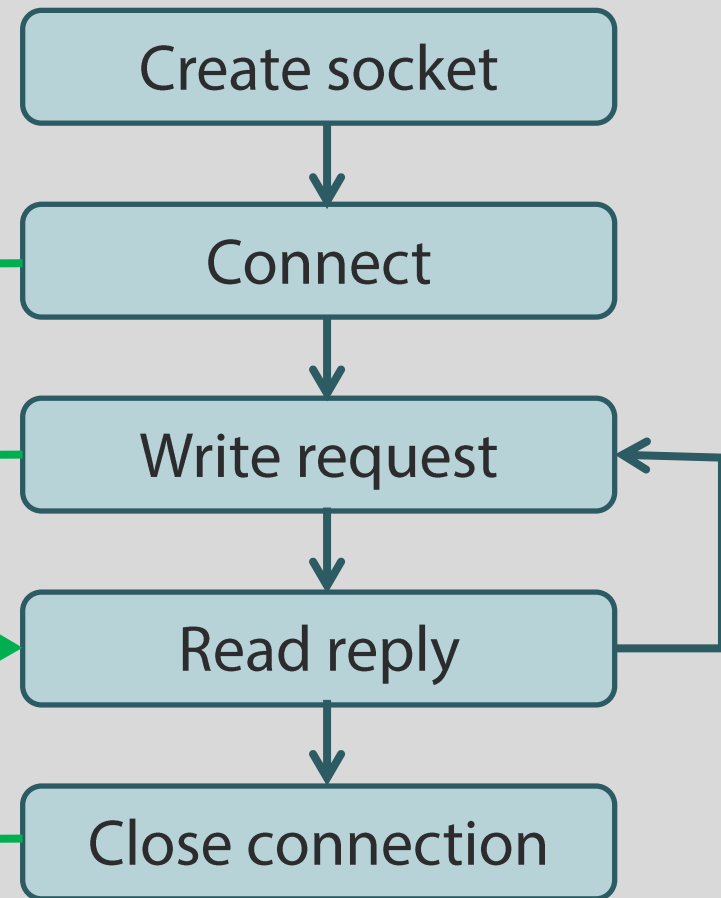
Finding the service
(the new way)


Doing it in Python

Server



Client





Association: {client IP, client port, server IP, server port} =
{10.1.2.17, 23451, 10.1.2.34, 80}

Creating a Socket

Transport type:
SOCK_STREAM
SOCK_DGRAM



```
sock = socket(domain, type, protocol)
```



Returns an integer
socket descriptor
or -1 on error



The address family. One of:
AF_UNIX
AF_INET
AF_INET6



Usually 0, system
selects protocol
based on domain
and type

Making the Connection

The socket
descriptor


The server's endpoint address
(sockaddr_in structure)

```
connect(sock, (struct sockaddr*)&server,  
        sizeof server)
```


Returns 0 on success
or -1 on error

After a successful connect(), sock
behaves as a file descriptor
referencing the connection
to the server

`/etc/hosts`



`gethostbyname(host)`

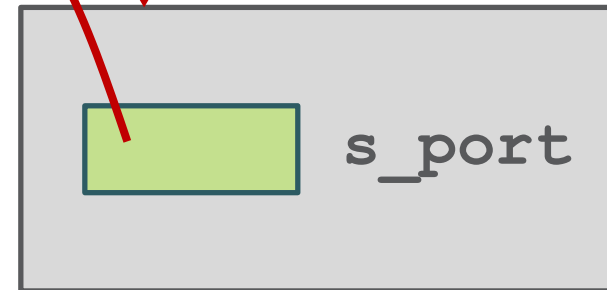


`struct hostent`


`/etc/services`





`getservbyname(service, protocol)`



`struct servent`



`struct sockaddr_in`



The /etc/services File

ftp-data	20/tcp
ftp	21/tcp
ssh	22/tcp
ssh	22/udp
telnet	23/tcp
smtp	25/tcp
domain	53/tcp
domain	53/udp
tftp	69/udp
http	80/tcp

Maps service names to
port number and protocol

For official port assignments see:

<http://www.iana.org/assignments/port-numbers>

Finding the Service

Service name

Protocol

`getservbyname("http", "tcp");`

`struct servent`


A null pointer is returned if the service name is not found

`getservbyport(80);`

The servent Structure

```
struct servent {  
    char    *s_name;           /* official service name */  
    char    **s_aliases;       /* alias list */  
    int      s_port;           /* port number */  
    char    *s_proto;          /* protocol */  
}
```

In network
byte order




Finding the Host

Host name



```
gethostbyname("venus.example.com");
```



struct hostent

A null pointer is returned
if the host name is not found

```
gethostbyname("192.168.1.44");
```

The hostent Structure

```
struct hostent {  
    char    *h_name;           /* official name of host */  
    char    **h_aliases;       /* alias list */  
    int      h_addrtype;       /* host address type */  
    int      h_length;         /* length of address */  
    char    **h_addr_list;     /* list of addresses */  
}  
/* For backward compatibility: */  
#define h_addr h_addr_list[0]
```

← "Canonical" name

← AF_INET or AF_INET6

← Null-terminated array of pointers to IP addresses


← Easy way to access the first address in the list

Putting it all Together

```
struct hostent      *host_info;  
struct servent      *serv_info;  
struct sockaddr_in  server;  
  
host_info = gethostbyname("venus.example.com");  
serv_info = getservbyname("http", "tcp");  
  
server.sin_family = AF_INET;  
memcpy(&server.sin_addr, host_info->h_addr, host_info->h_length);  
server.sin_port = serv_info->s_port;  
connect(sock, (struct sockaddr*)&server, sizeof server)
```

Demonstration

Old-fashioned TCP client



Protocol Independence

Traditional API
(`gethostbyname()`, etc)
makes it hard to be
"protocol independent"

Growing need for
IPv4 / IPv6 operability

`getaddrinfo()`
more complex API ...
but easier to achieve
protocol independence

Also ...
traditional API is not
suitable for multi-threaded
applications

Finding the Service (the Modern Way)

Machine name
↓
`getaddrinfo ("venus.example.com",`
"http", ← Service Name
&hints, ← Selection criteria
&result);
↓
Linked list of endpoint addresses

Returns zero on success, non-zero on error

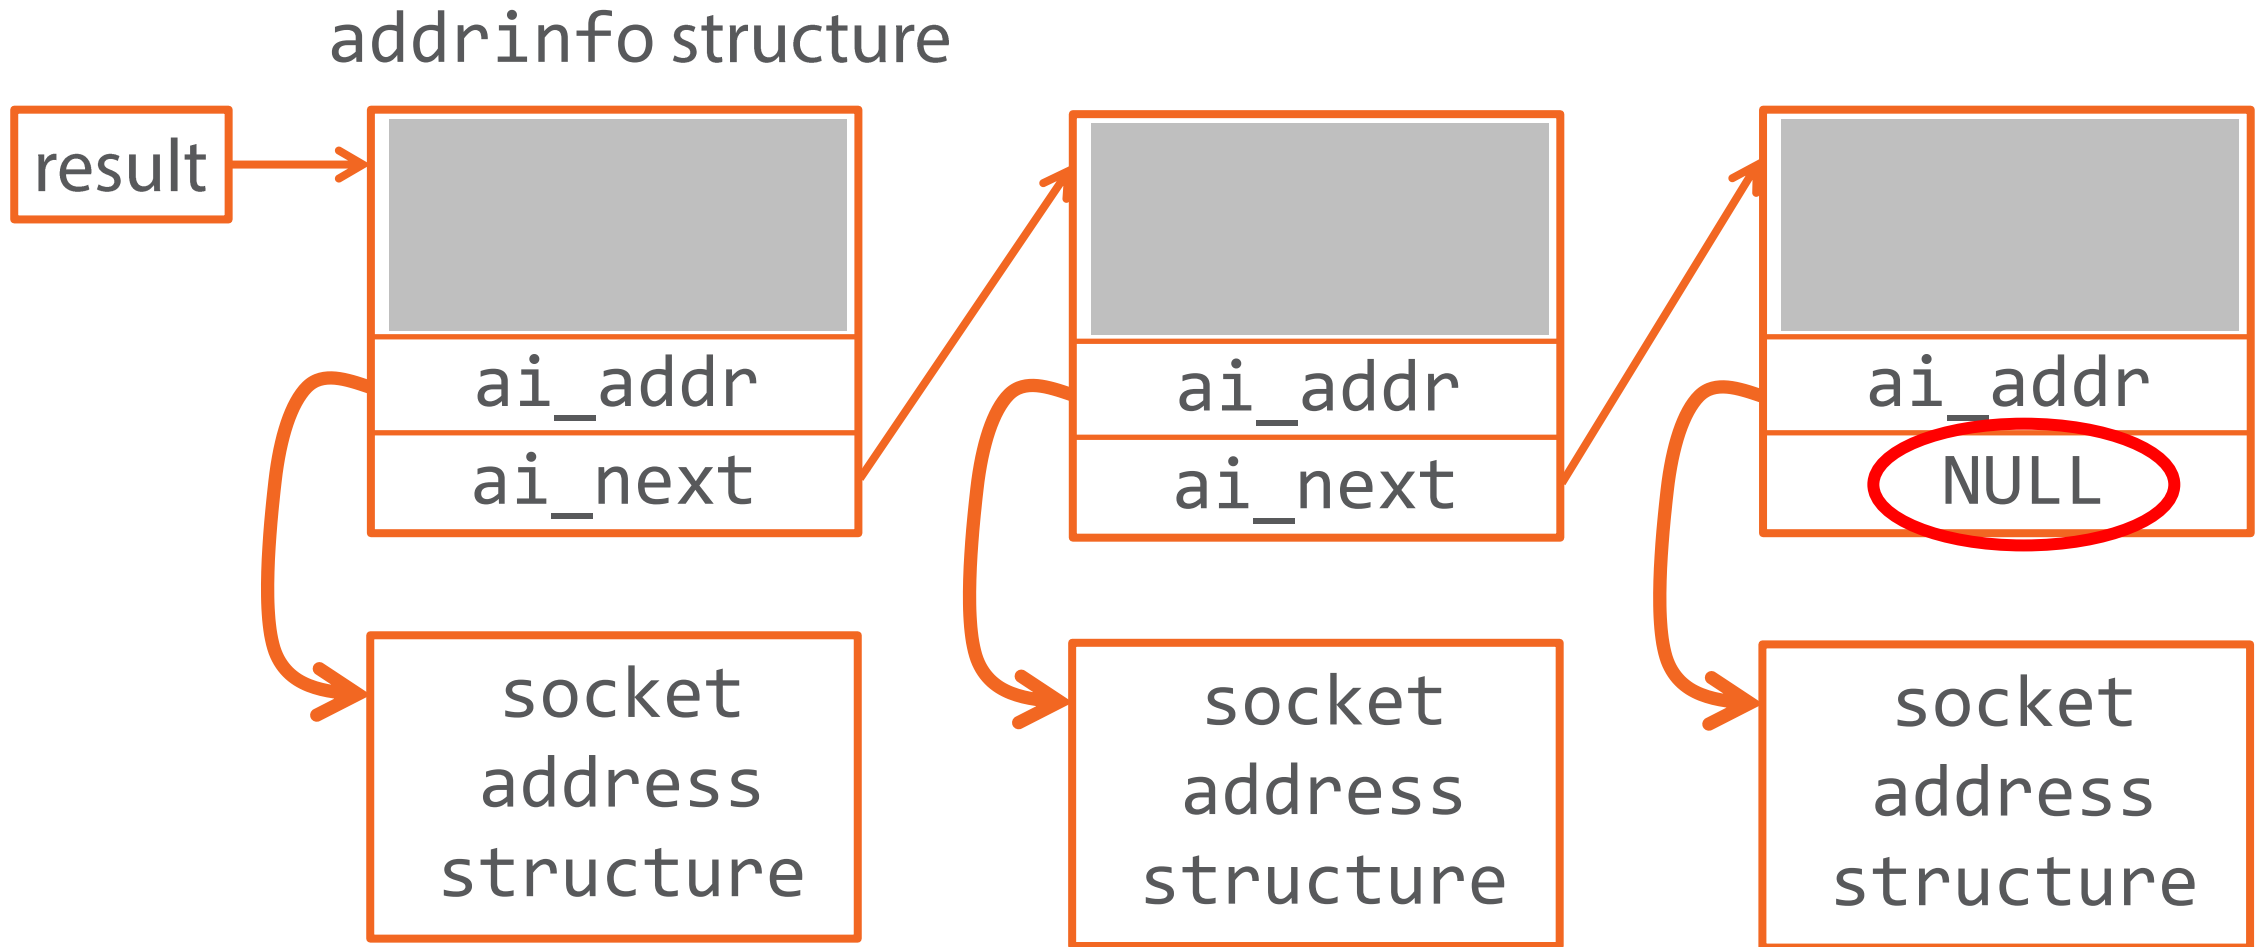
Providing Hints

The `hints` argument of `getaddrinfo()` lets you constrain the types of endpoint addresses you want.

For example, to return only IPV6 UDP endpoints:

```
struct addrinfo  hints;  
  
memset(&hints, 0, sizeof(struct addrinfo));  
hints.ai_family = AF_INET6;  
hints.ai_socktype = SOCK_DGRAM
```

Data Structures




The addrinfo Structure

```
struct addrinfo {  
    int          ai_flags;  
    int          ai_family; ← AF_INET or AF_INET6  
    int          ai_socktype; ← SOCK_DGRAM or SOCK_STREAM  
    int          ai_protocol;  
    socklen_t    ai_addrlen; ← The length of the sockaddr  
    struct sockaddr *ai_addr; ← Pointer to socket address  
    char         *ai_canonname;  
    struct addrinfo *ai_next; ← Pointer to next structure  
                                in linked list  
};
```

Demonstration

Modern TCP client



Doing it in Python



```
#!/usr/bin/python3
# The core of a Python TCP client

from socket import *


s = socket(AF_INET, SOCK_STREAM)
port = getservbyname("rot13")
s.connect(("localhost", port))

s.send(...)
s.recv(...)
```

Must be an integer

Demonstration

Python TCP client



Module Summary


Finding the service

Finding the host

Connecting to the server

Protocol Independence

Doing it in Python



Moving Forward ...



Coming up in the next module:

UDP clients and servers

TFTP client

UDP broadcasting