

Multi-threaded Concurrency



Chris Brown



In This Module ...

Thread concepts

The pthreads API


Multi-threaded
servers and clients

Processor farms

Demonstration:
Counting primes

Writing "thread-
safe" code

Multi-threading Illustrated



Reasons for Multi-threading

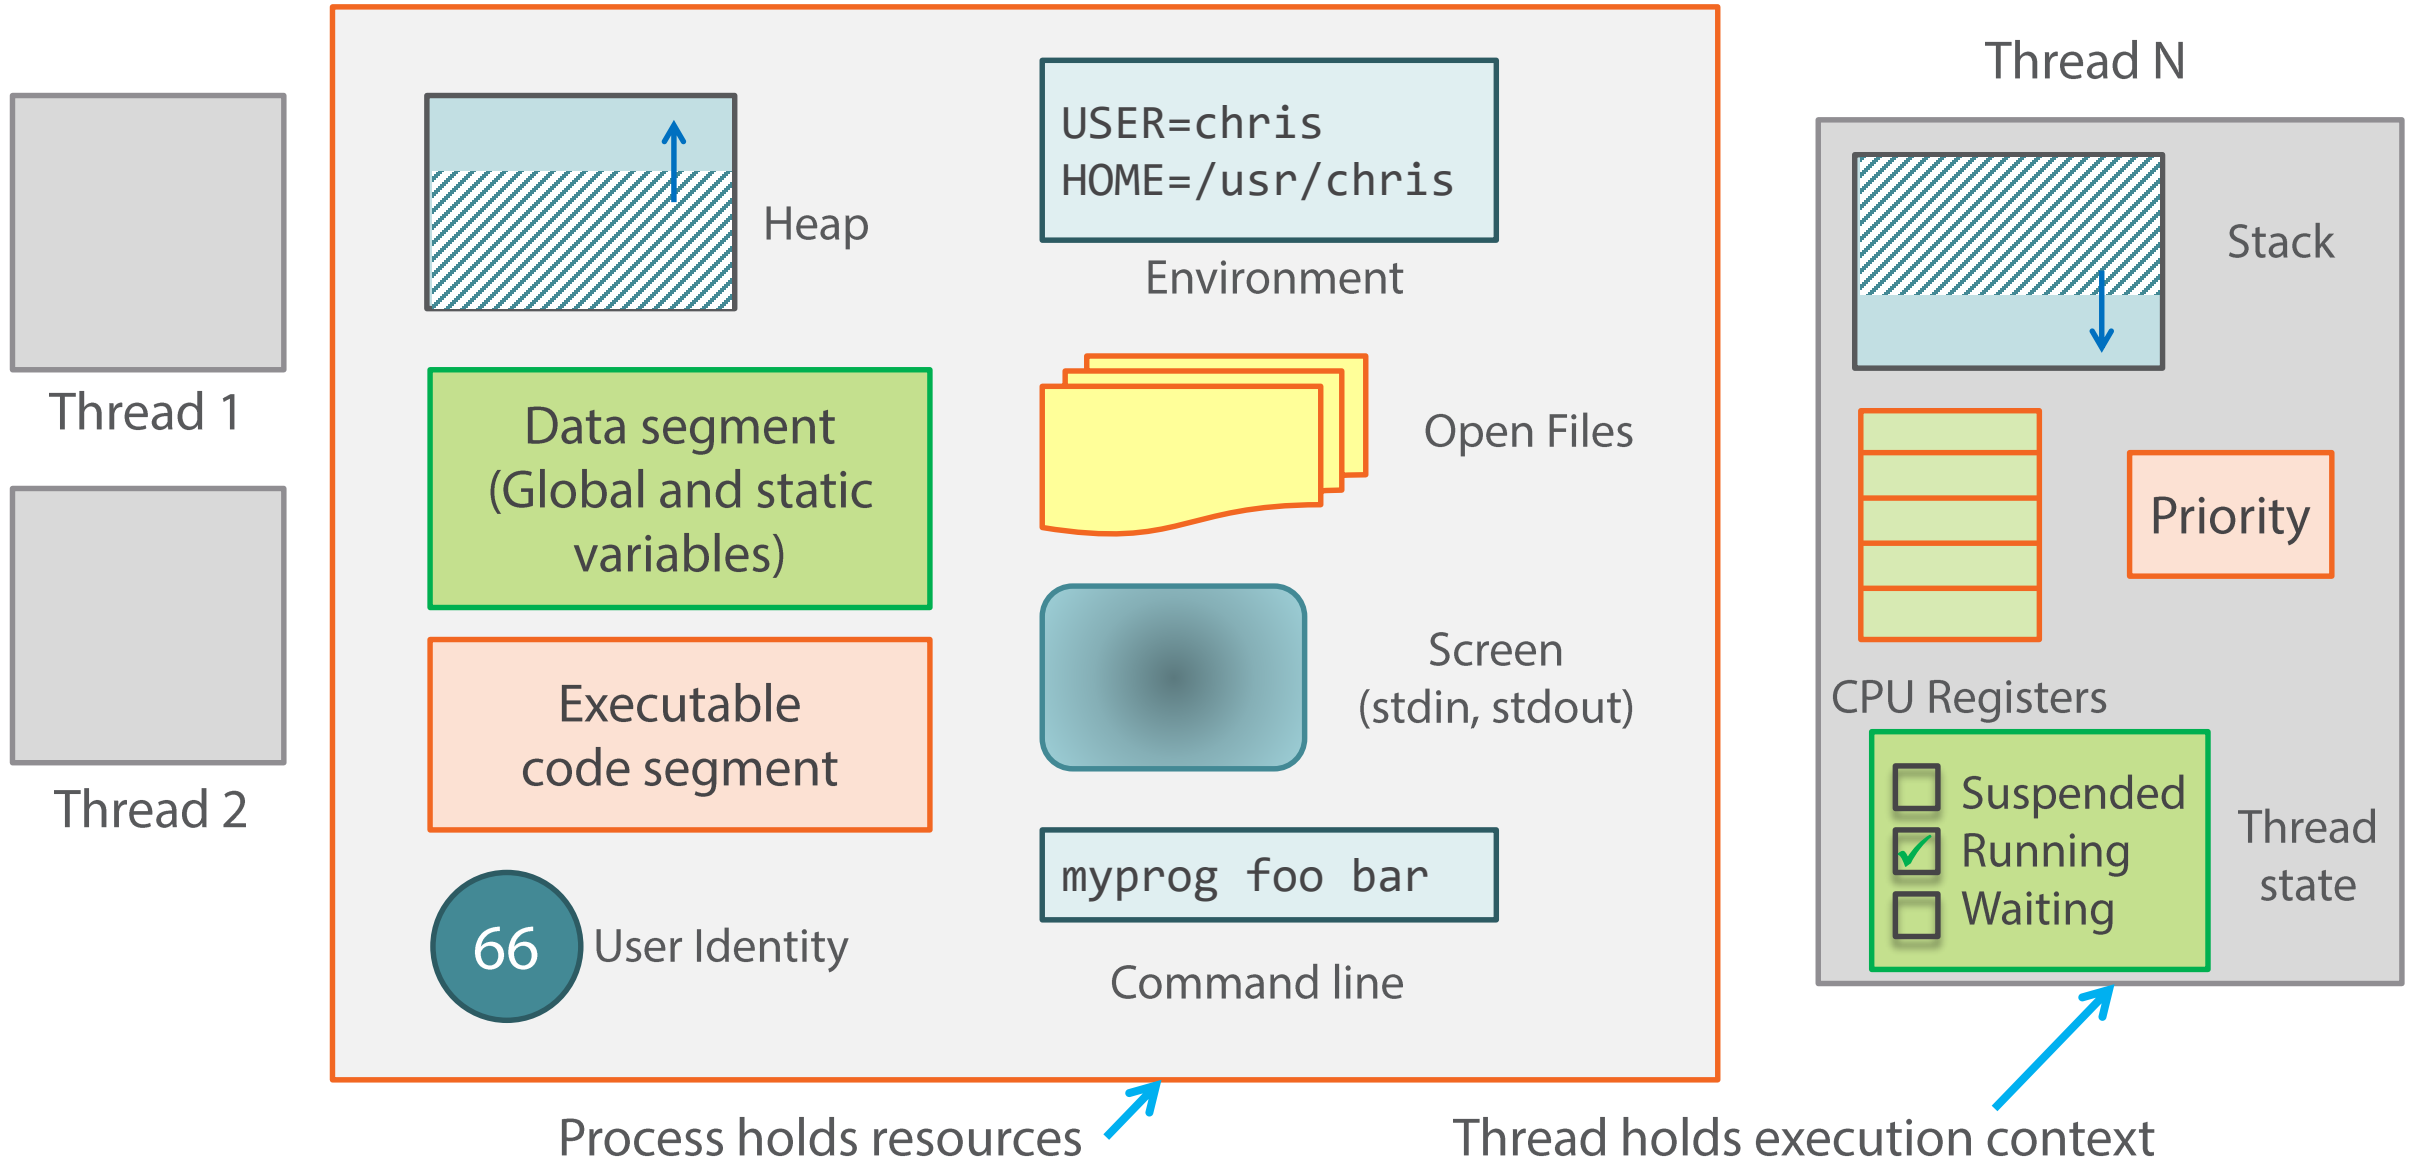
Expressing logical
concurrency

Implementing
background tasks

Concurrent servers

Exploiting multi-
processor hardware

Threads and Processes



Shared and Not Shared

Threads share:

- Code
- Global and static variables
- Open file descriptors

Threads do not share:

- Variables local to functions

Pthreads



POSIX 1003.1c



A standardised set of C library routines for thread creation and management

— Upwards of 60 functions

Thread Creation

Thread Attribute Object
(NULL for defaults)

`pthread_create(&handle, &attr, func, arg)`

Returns 0 if OK,
Nonzero if error

Thread handle
returned here

```
void *func(void *arg)
{
    //Thread function
}
```


Thread Termination

- A thread can terminate in several ways:
 1. By calling `pthread_exit(exit_status)`
 2. By returning from its top-level function
 3. By some other thread sending a cancellation request:
`pthread_cancel(handle);`
- Parent can wait for thread to finish and get exit status:
 - `pthread_join(handle, &exit_status);`
- Parent should detach thread if they don't need to join on it:
`pthread_detach(handle);`

Thread Life Cycle

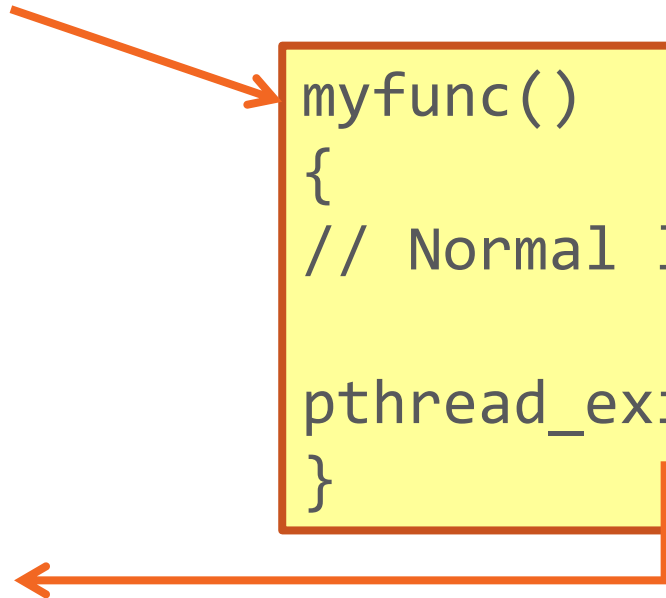
Parent Thread

Child Thread

```
pthread_create(..., myfunc, ...);
```

```
myfunc()  
{  
    // Normal life of thread  
    pthread_exit(status);  
}
```

```
pthread_join(handle, &status);
```



Demonstration

Simple thread example



Thread Example

```
#include <pthread.h>
#include <stdio.h>

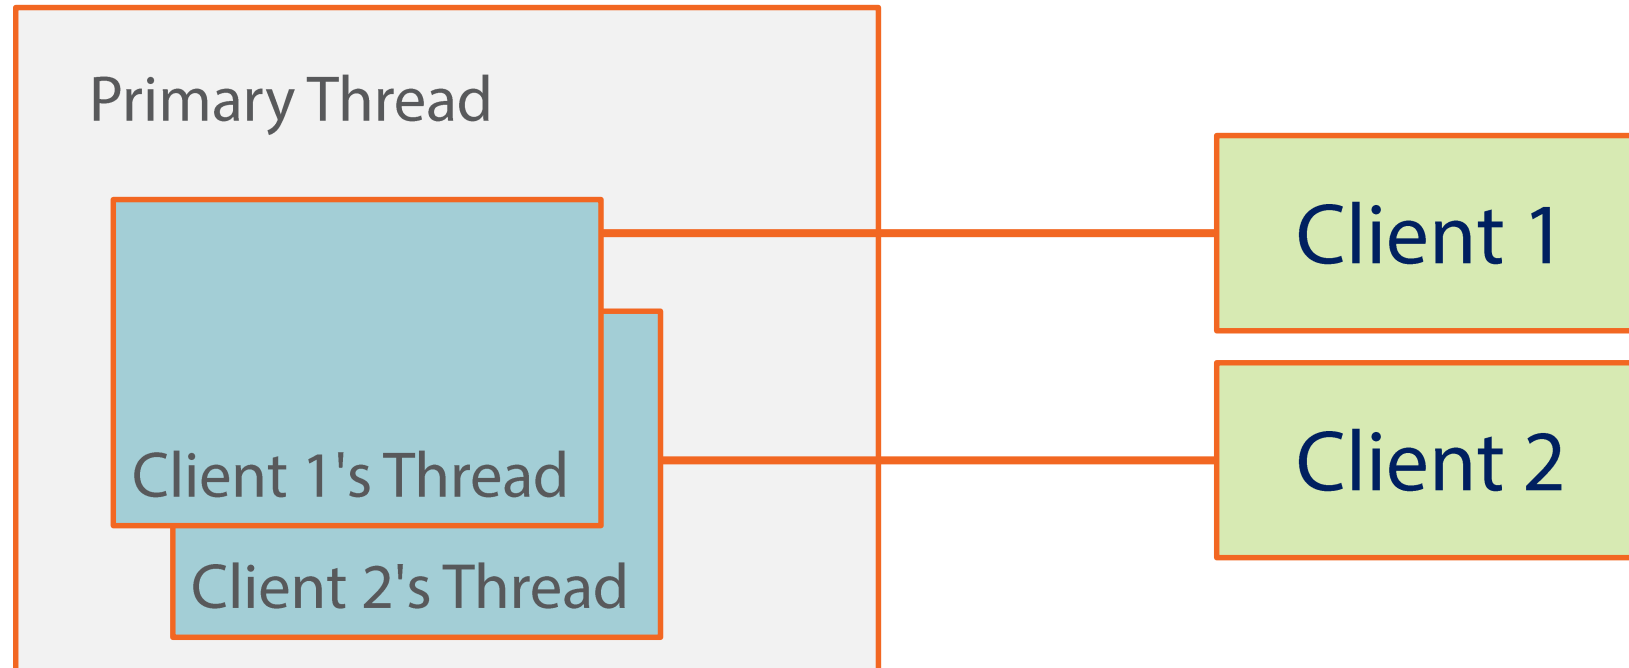
void *func(void *arg)
{
    printf("child thread says %s\n", (char *)arg);
    pthread_exit((void *)99);
}

int main()
{
    pthread_t handle;
    int exitcode;


    pthread_create(&handle, NULL, func, "hi!");
    printf("primary thread says hello\n");
    pthread_join(handle, (void **)&exitcode);
    printf("exit code %d\n", exitcode);
}
```

Multi-threaded Server


- Thread-per-client is an elegant model for concurrent servers
 - Efficient
 - Easy to keep per-client state (local variables)
 - Easy to share state (global variables)



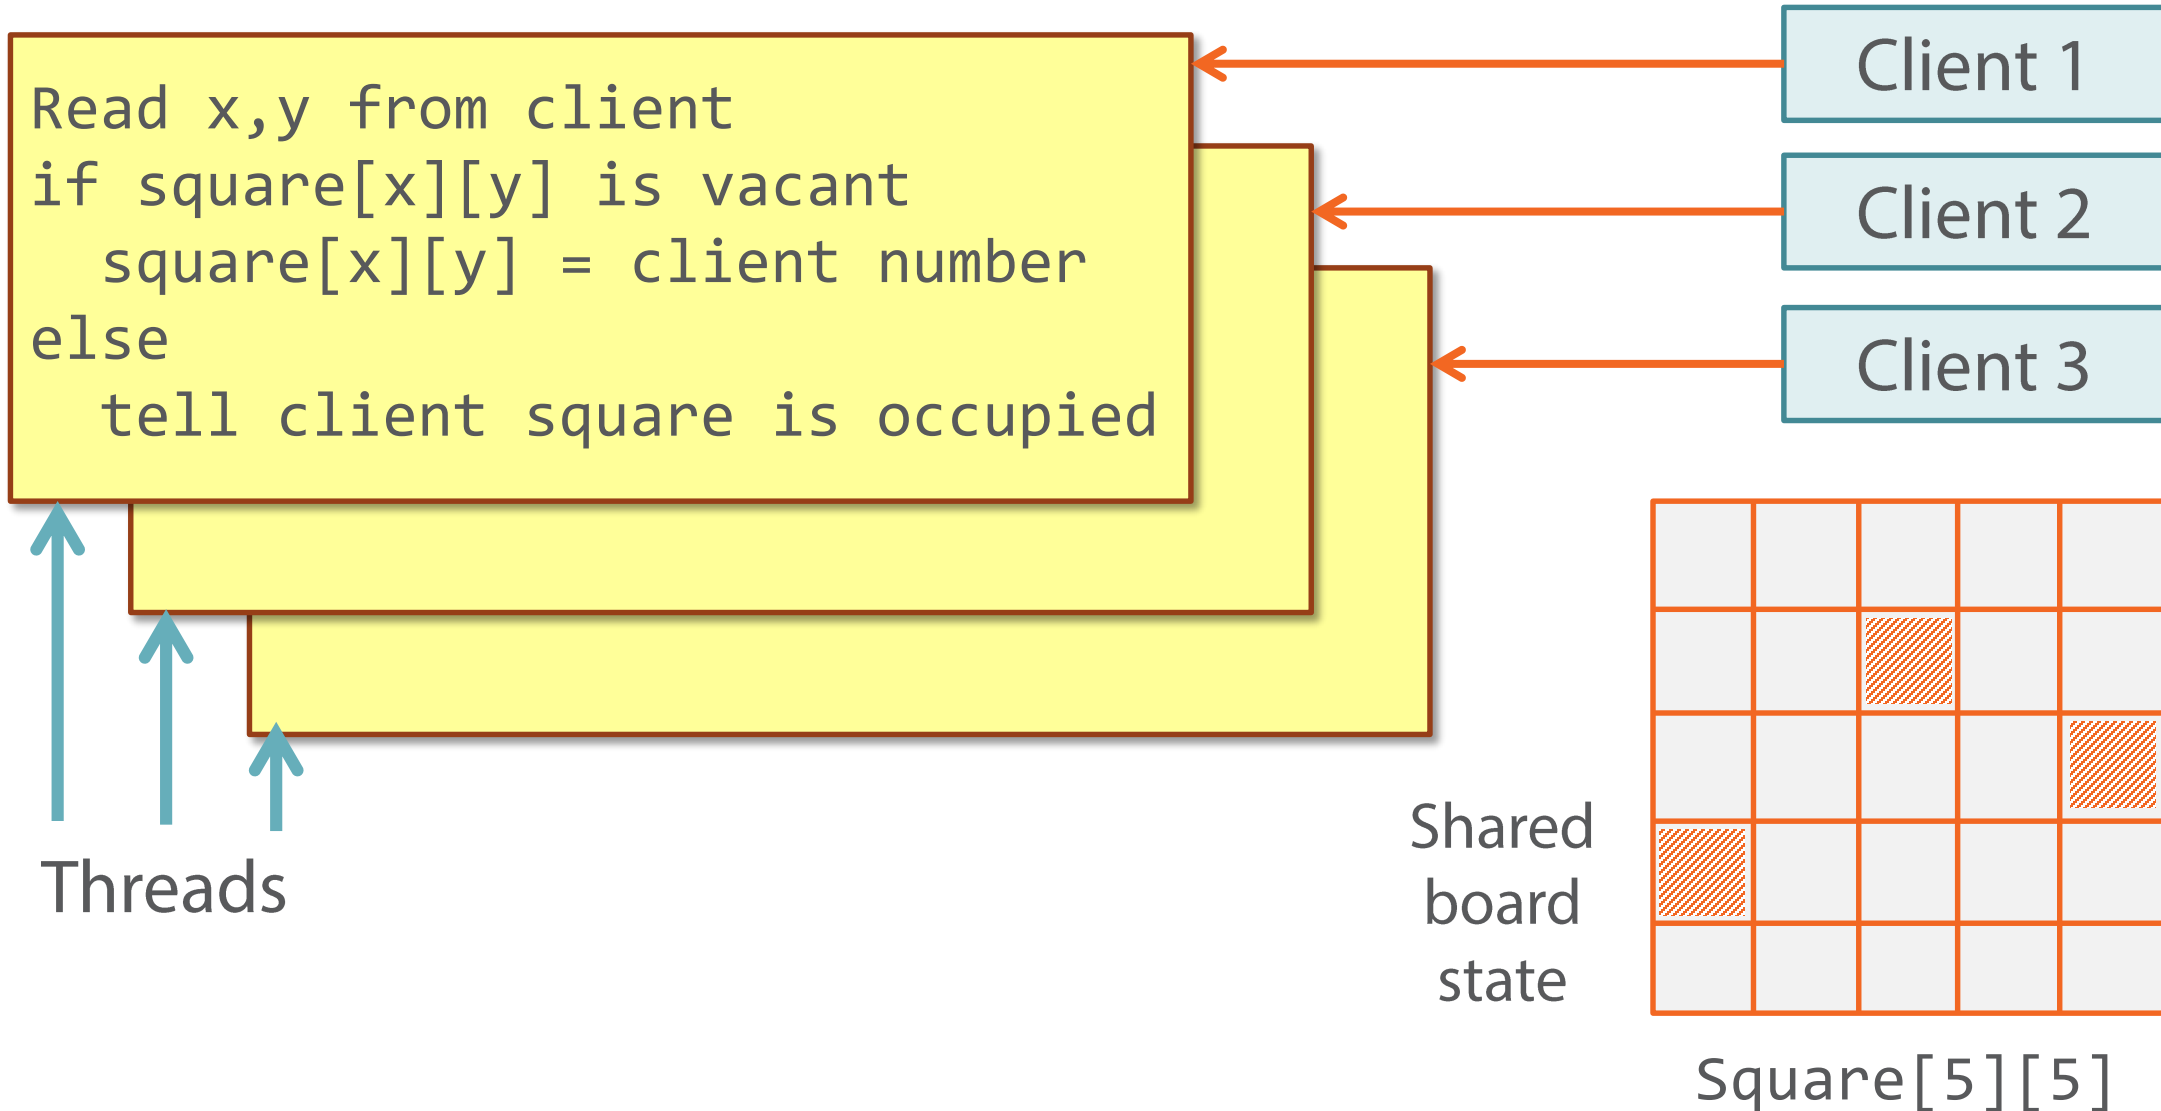
Concurrent Server




Processor Farm



Sharing Client State



Thread Safety



Problems can arise when multiple asynchronous threads access shared data

Mutual exclusion locks ("mutexes") control entry to code sections that update or access shared state

"Thread-safe" code



Library functions you call must also be thread-safe

Pthread Mutexes

```
#include <pthread.h>

static int shared_data;
pthread_mutex_t mylock = PTHREAD_MUTEX_INITIALIZER;

func( ... )
{
    pthread_mutex_lock(&mylock);
    // Update or access shared_data here
    pthread_mutex_unlock(&mylock);
}
```

Demonstration

Why do we need
mutexes?



Processor Farms

A common model for decomposing computationally-intensive tasks

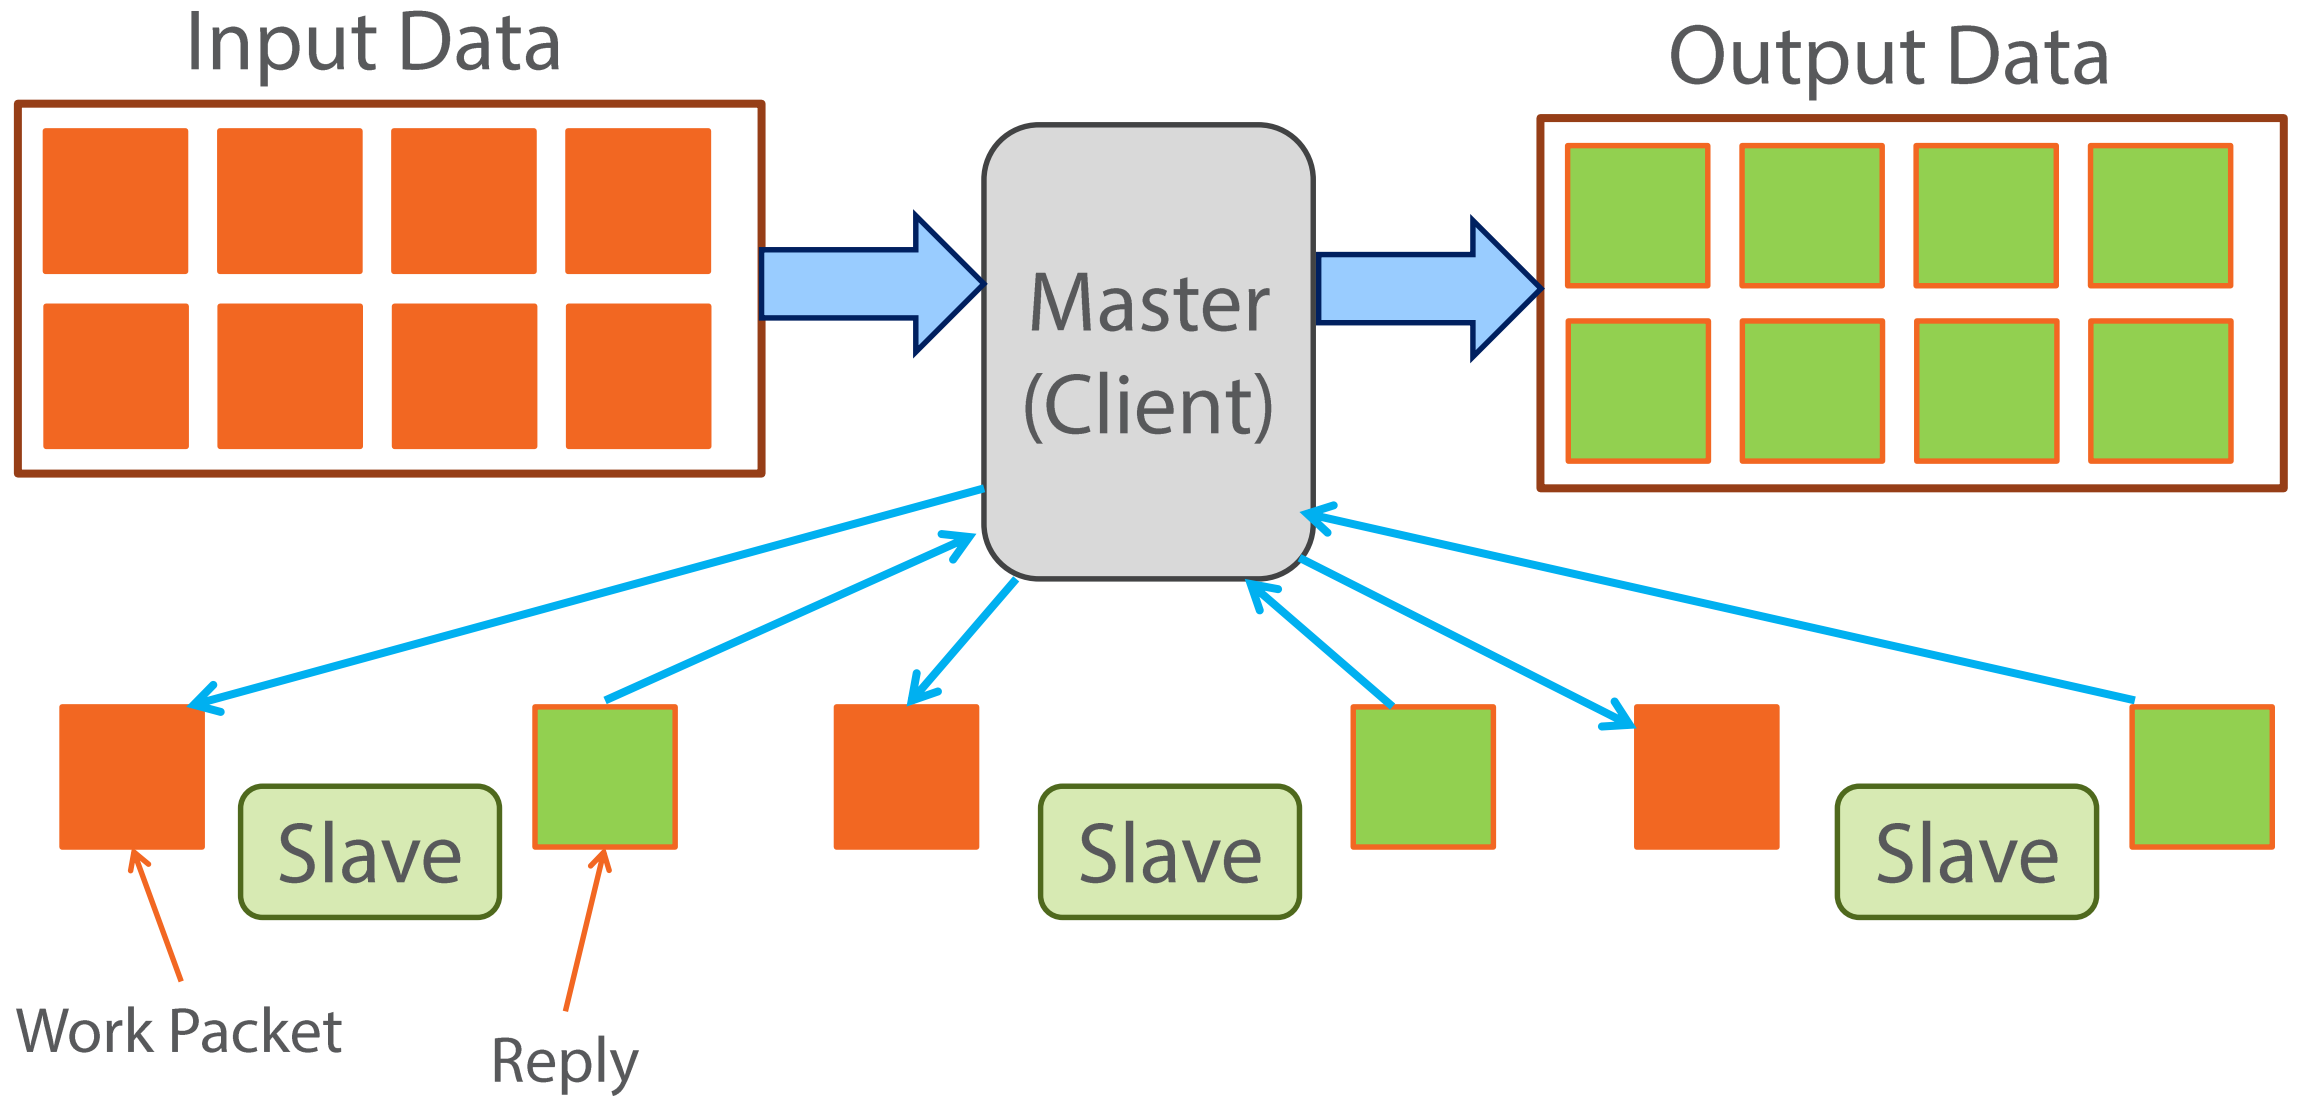
Input data set broken down into "work packets"
— each processed independently of the others



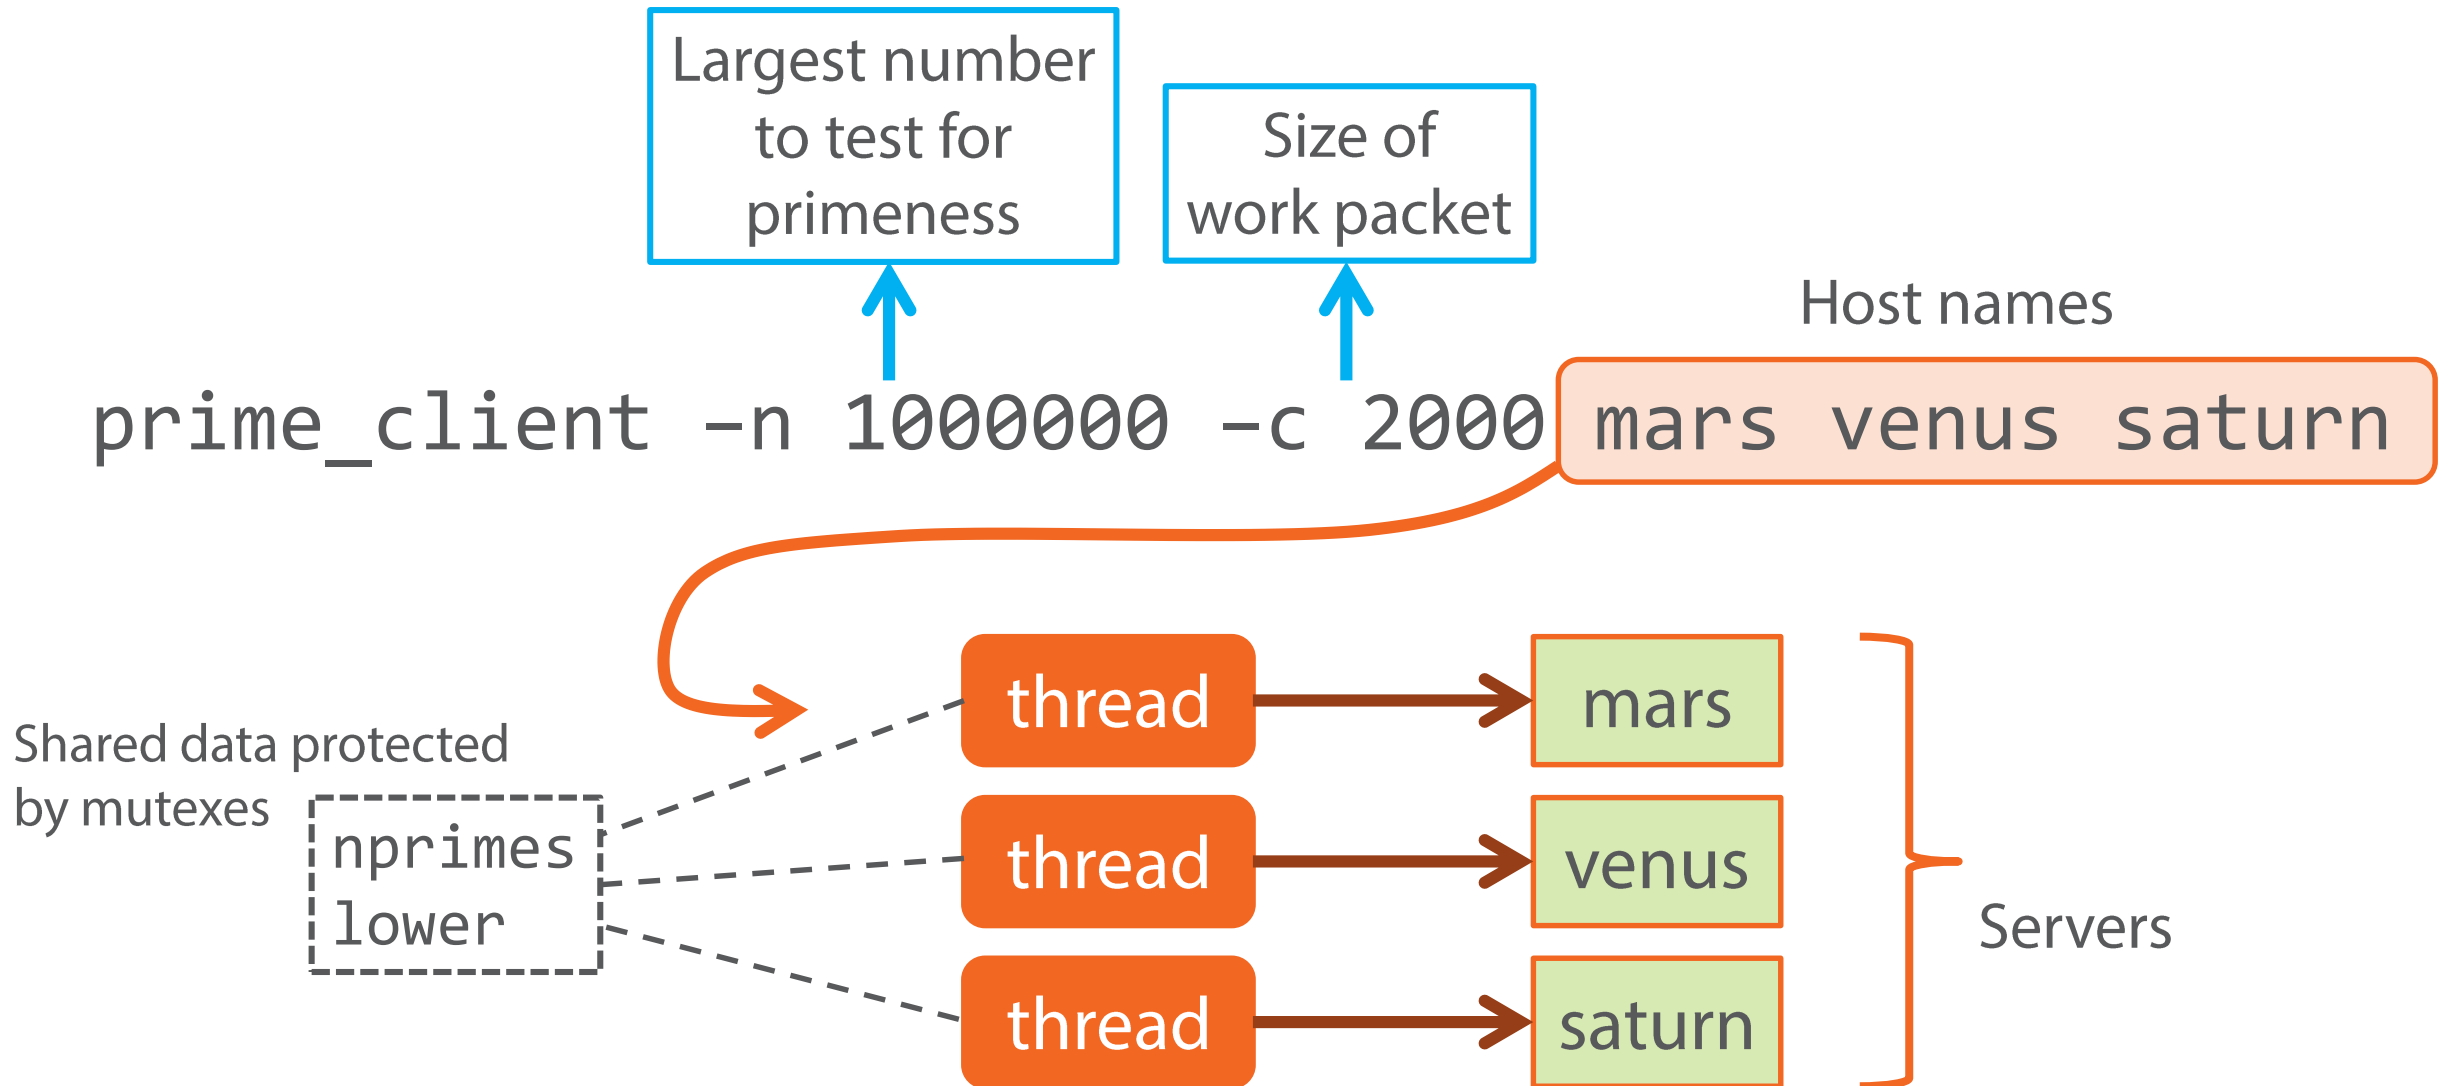
Need $\approx 10\times$ as many work packets as processors
— Automatic load balancing

Goal: Linear speedup

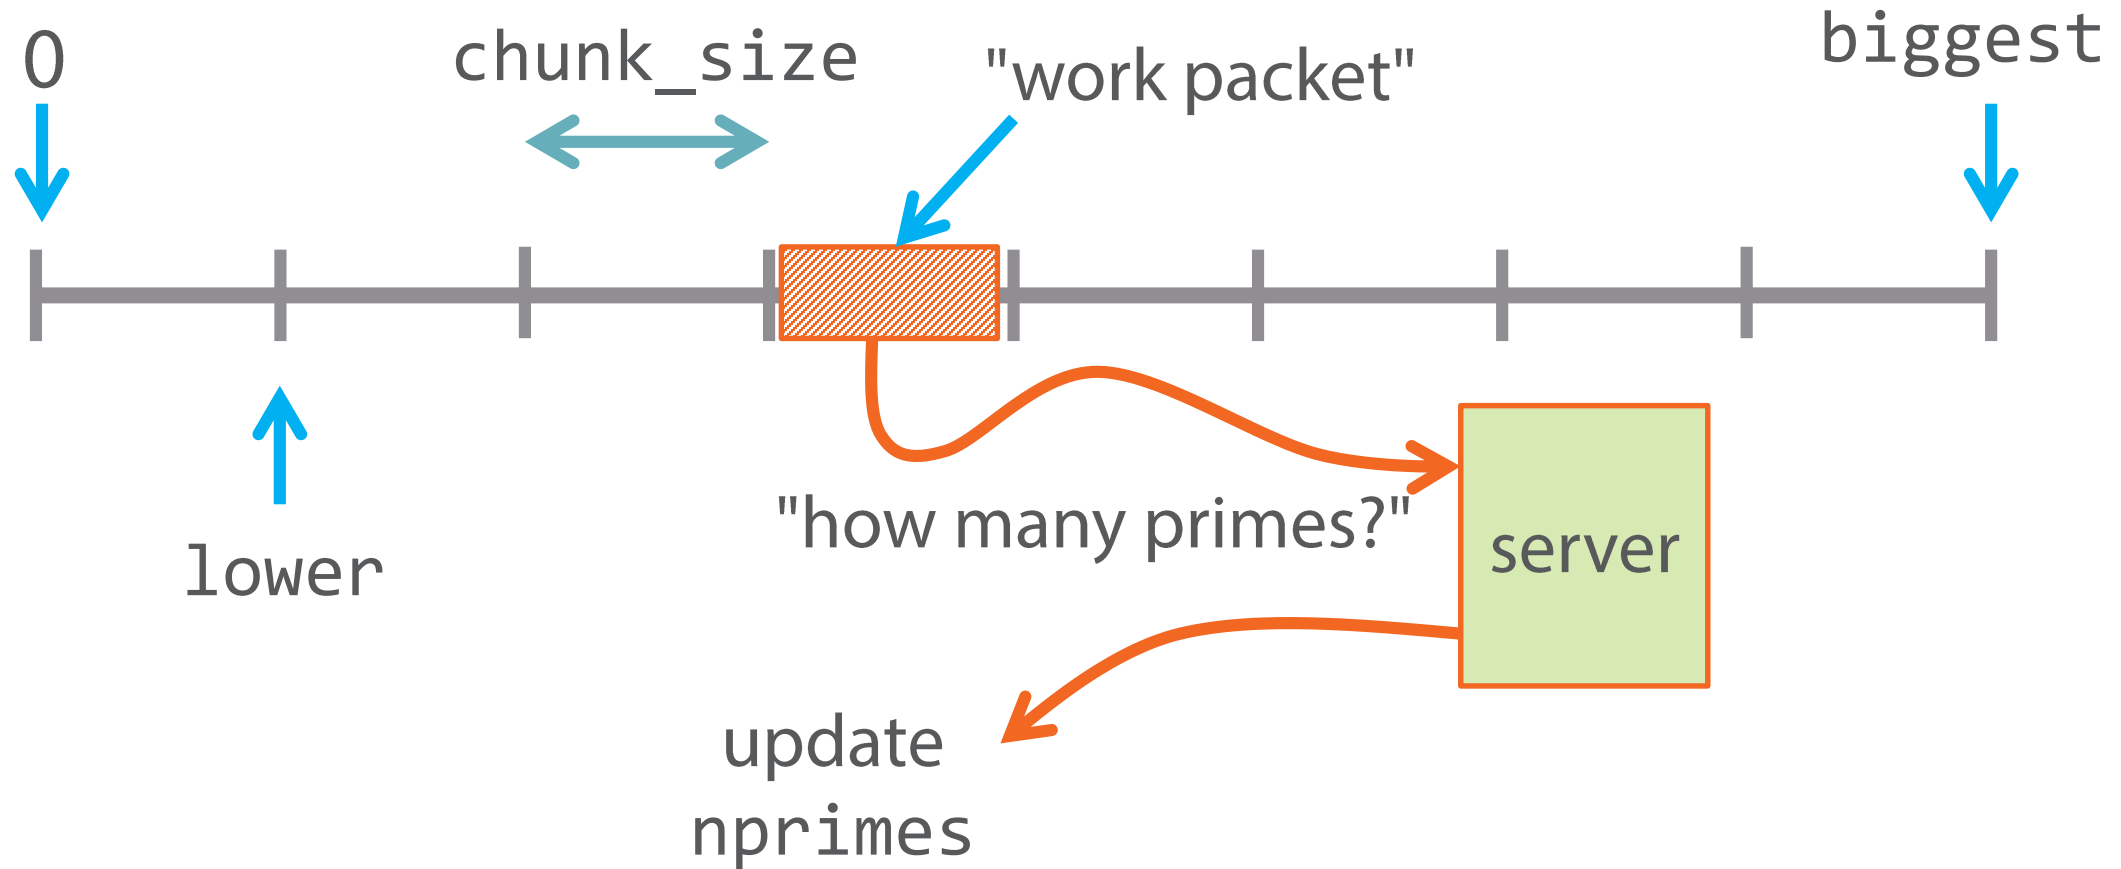
Processor Farm Illustrated



A Processor Farm to Count Prime Numbers




Dividing up the Work



Demonstration

Processor Farm



Module Summary



In this module:

Threads compared to processes

The pthreads API

Thread-safe code: accessing shared data

The processor farm model

A processor farm to count prime numbers

Course Summary



Congratulations!

In this course:

The characteristics of TCP and UDP protocols

Writing TCP-based servers

Writing TCP-based clients

UDP servers and clients

Writing concurrent servers

Writing concurrent clients using threads